


Introduction to Distributed Systems and Characterisation

Dr. Rajkumar Buyya



Cloud Computing and **D**istributed **S**ystems (CLOUDS) Laboratory
School of Computing and Information Systems
The University of Melbourne, Australia

<http://www.buyya.com>

<http://clouds.cis.unimelb.edu.au/~rbuyya/>

Most concepts are
drawn from Chapter 1


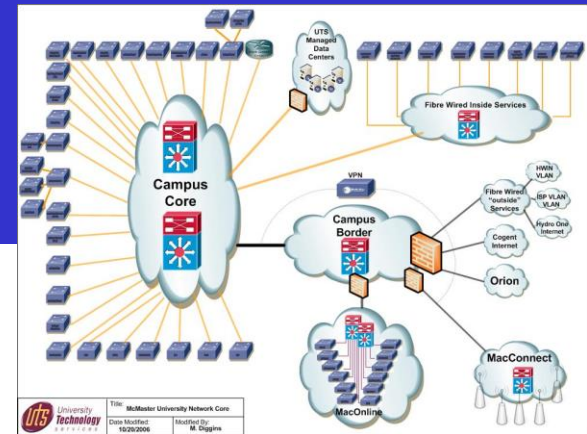
Presentation Outline

- Introduction
- Defining Distributed Systems
- Characteristics of Distributed Systems
- Example Distributed Systems
- Challenges of Distributed Systems
- Summary

Introduction

- Networks of computers are everywhere!
 - Mobile phone networks
 - Corporate networks
 - Factory networks
 - Campus networks
 - In-car networks
 - **Internet of Things (IoT)**
 - On board networks in planes and trains
- This subject aims:
 - to cover characteristics of networked/distributed computing systems and applications
 - to present the main concepts and techniques that have been developed to help in the tasks of designing and implementing systems and applications that are based on networks.

How mobile networks work




Defining Distributed Systems

- *“A system in which hardware or software components located at **networked** computers communicate and coordinate their actions only by **message passing**.” [Coulouris]*
- *“A distributed system is a collection of **independent** computers **that appear** to the users of the system as a single computer.” [Tanenbaum]*
- Example Distributed Systems:
 - Cluster:
 - *“A type of parallel or distributed processing system, which consists of a collection of interconnected **stand-alone** computers cooperatively **working together** as a single, integrated computing resource” [Buyya].*
 - Cloud:
 - *“a type of parallel and distributed system consisting of a collection of **interconnected** and **virtualised computers** that are **dynamically provisioned** and presented as one or more unified computing resources based on **service-level agreements** established through negotiation between the service provider and consumers” [Buyya].*


Leslie Lamport's Definition

- *"A distributed system is one on which I **cannot** get any work done because some machine I have never heard of has crashed."*
- Leslie Lamport – a famous researcher on timing, message ordering, and clock synchronization in distributed systems.



Networks vs. Distributed Systems

- **Networks:** A media for interconnecting local and wide area computers and exchange messages based on protocols. Network entities are visible and they are explicitly addressed (IP address).
- **Distributed System:** existence of multiple autonomous computers is transparent
- However,
 - many problems (e.g., openness, reliability) in common, but at different levels.
 - Networks focuses on packets, routing, etc., whereas distributed systems focus on applications.
 - Every distributed system relies on services provided by a computer network.



Reasons for Distributed Systems

■ Functional Separation:

- Existence of computers with different capabilities and purposes:
 - Clients and Servers
 - Data collection and data processing

■ Inherent distribution:

- Information:
 - Different information is created and maintained by different people (e.g., Web pages)
- People
 - Computer supported collaborative work (virtual teams, engineering, virtual surgery)
- Retail store and inventory systems for supermarket chains (e.g., Coles, Woolworths)

■ Power imbalance and load variation:


- Distribute computational load among different computers.

■ Reliability:

- Long term preservation and data backup (replication) at different locations.


■ Economies:

- Sharing a printer by many users and reduce the cost of ownership.
- Building a supercomputer out of a network of computers.



Consequences of Distributed Systems

- Computers in distributed systems may be on separate continents, in the same building, or the same room. DSs have the following consequences:
 - Concurrency – each system is autonomous.
 - Carry out tasks independently
 - Tasks coordinate their actions by exchanging messages.
 - Heterogeneity
 - No global clock
 - Independent Failures



Characteristics of Distributed Systems

- Parallel activities

- Autonomous components executing concurrent tasks

- Communication via message passing

- No shared memory

- Resource sharing


- Printer, database, other services

- No global state

- No single process can have knowledge of the current global state of the system



- No global clock

- Only limited precision for processes to synchronize their clocks



Goals of Distributed Systems

- Connecting Users and Resources
- Transparency
- Openness
- Scalability
- Enhanced Availability



Examples of Distributed Systems

- They (DS) are based on familiar and widely used computer networks:

- Internet
- Intranets, and
- Wireless networks



- Example DS and its Applications:

- Web (and many of its applications like Online bookshop)
- Data Centers and Clouds
- Wide area storage systems
- Banking Systems
- User-level communication (Facebook, Zoom)

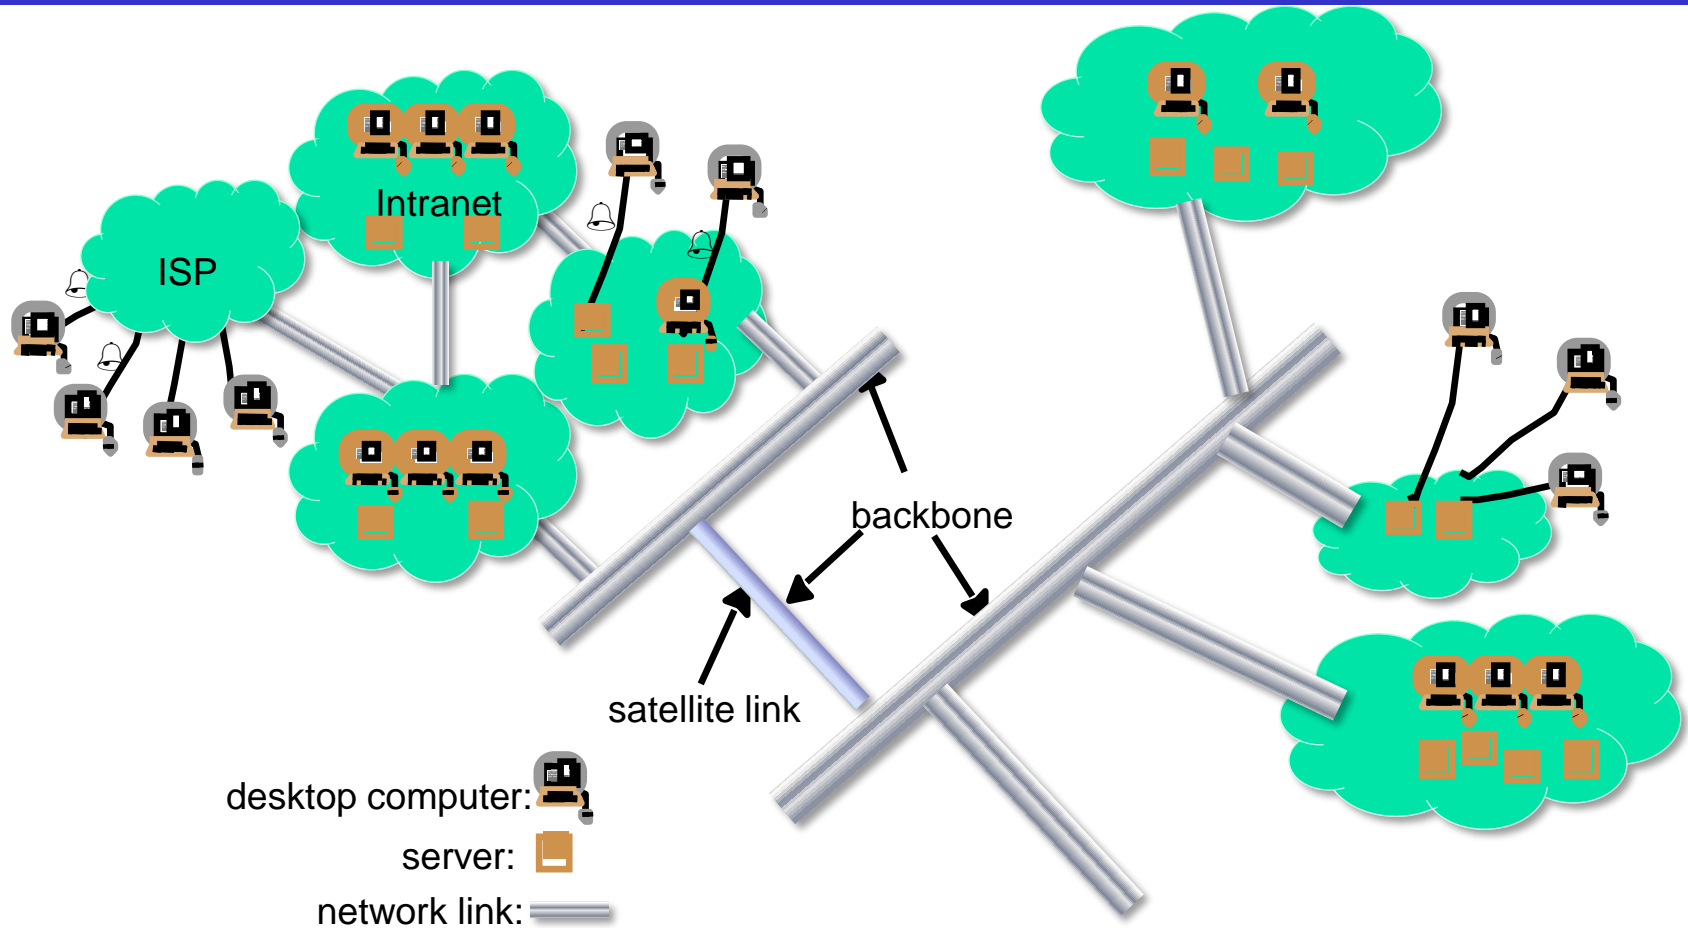


Selected application domains and associated networked applications

<i>Finance and Commerce</i>	eCommerce e.g. Amazon and eBay , PayPal , online banking and trading
<i>The Information Society</i>	Web information and search engines, ebooks, Wikipedia; social networking: Facebook , Twitter , and WeChat .
<i>Creative Industries and Entertainment</i>	Online gaming, music and film in the home, user-generated content, e.g. YouTube , Flickr , Netflix
<i>Healthcare</i>	Health informatics, on online patient records, monitoring patients (Metro South Health hospital trial in Queensland)
<i>Education</i>	e-learning, virtual learning environments; distance learning. e.g., Coursera
<i>Transport and Logistics</i>	GPS in route finding systems, map services: Google Maps , Google Earth
<i>Science and Engineering</i>	Cloud computing as an enabling technology for collaboration between scientists (LHC , LIGO)
<i>Environmental Management</i>	Sensor networks to monitor earthquakes, floods or tsunamis (Bureau of Meteorology flood warning system)

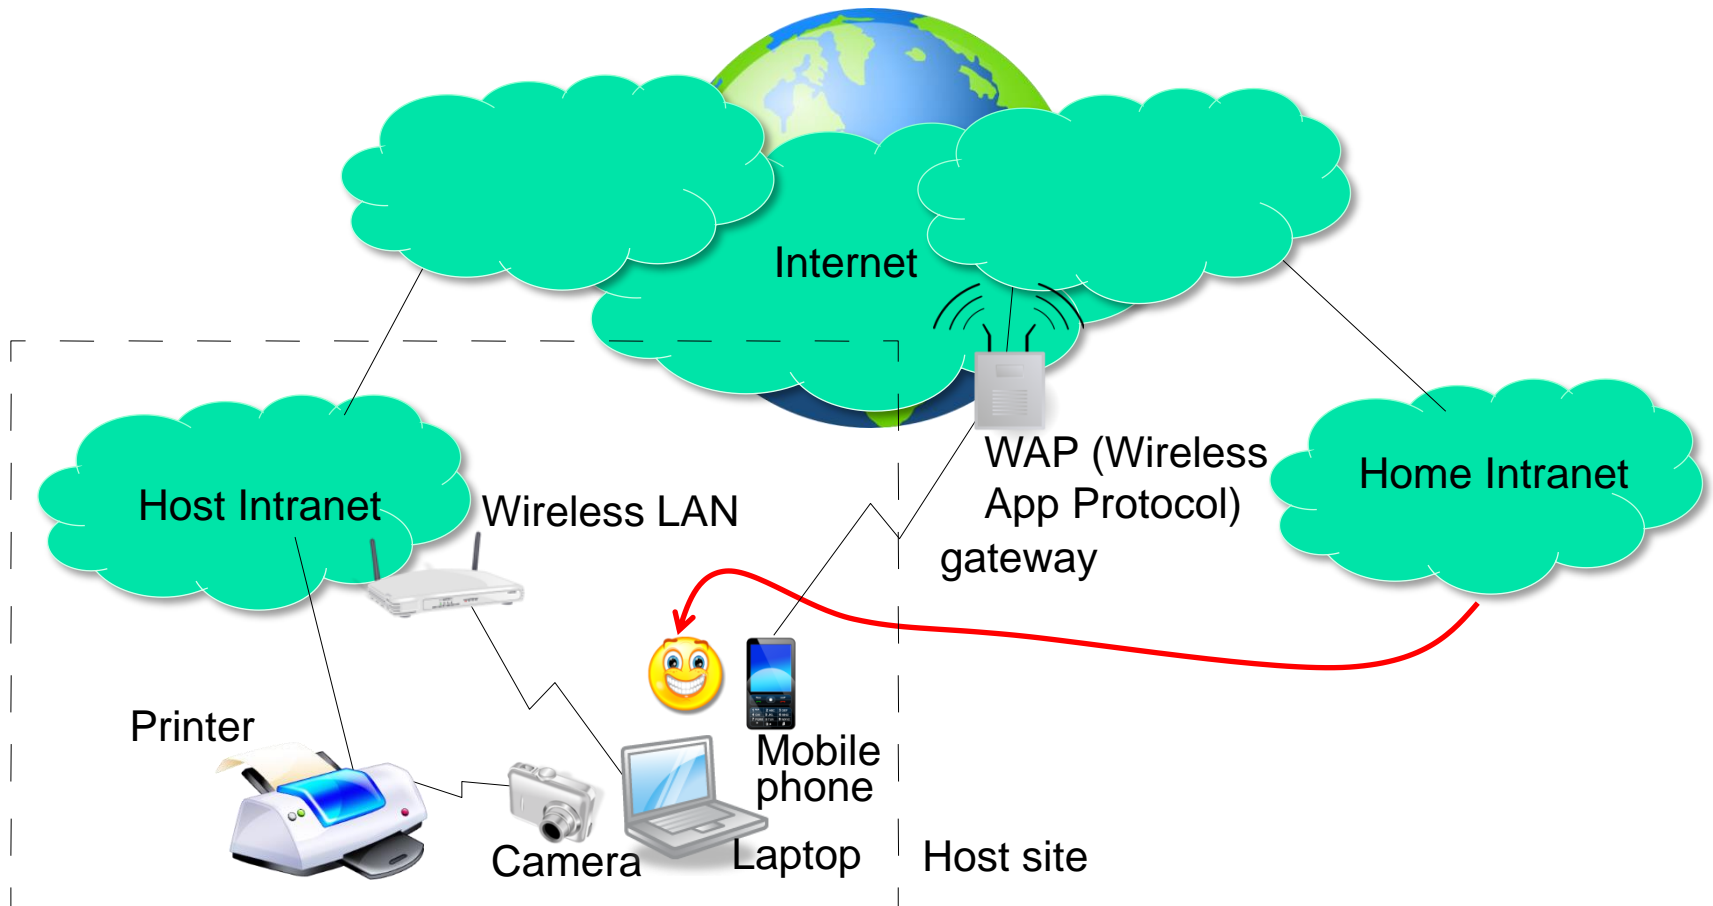
A typical portion of the Internet and its services:

Multimedia services providing access to music, radio, TV channels, and video conferencing supporting several users.




- The Internet is a vast collection of computer networks of many different types and hosts various types of services.

Mobile and ubiquitous computing: portable and handheld devices in a distributed system




- Supports continued access to Home intranet resources via wireless and provision to utilise resources (e.g., printers) that are conveniently located (location-aware computing).

Resource sharing and the Web: open protocols, scalable servers, and pluggable browsers



Business Example and Challenges

- Online bookstore (e.g. in World Wide Web)
 - Customers can connect their computer to your computer (web server):
 - Browse your inventory
 - Place orders
 - ...



This example has been adapted from **Torbin Weis**, Berlin University of Technology

Business Example – Challenges I

■ What if

- Your customer uses a completely different hardware? (PC, MAC, iPad, Mobile...)
- ... a different operating system? (Windows, Unix,...)
- ... a different way of representing data? (ASCII, EBCDIC,...)
- **Heterogeneity**

■ Or

- You want to move your business and computers to the Caribbean (because of the **weather** or **low tax**)?
- Your client moves to the Caribbean (more likely)?
- **Distribution transparency**

Business Example – Challenges II

■ What if

- Two customers want to order the same item at the same time?
- **Concurrency**

■ Or

- The database with your inventory information crashes?
- Your customer's computer crashes in the middle of an order?
- **Fault tolerance**

Business Example – Challenges III

■ What if


- Someone tries to break into your system to steal data?
- ... sniffs for information?
- ... your customer orders something and doesn't accept the delivery saying he didn't?

■ **Security**

■ Or

- You are so successful that millions of people are visiting your online store at the same time?

■ **Scalability**



Business Example – Challenges IV

- When building the system...
 - Do you want to write the whole software on your own (network, database,...)?
 - What about updates, new technologies?
 - **Reuse** and **Openness** (Standards)



Overview Challenges I

- Heterogeneity

- Heterogeneous components must be able to interoperate

- Distribution transparency

- Distribution should be hidden from the user as much as possible

- Fault tolerance

- Failure of a component (partial failure) should not result in failure of the whole system

- Scalability

- System should work efficiently with an increasing number of users
- System performance should increase with inclusion of additional resources

Overview Challenges II

- Concurrency

- Shared access to resources must be possible

- Openness


- Interfaces should be publicly available to ease inclusion of new components

- Security

- The system should only be used in the way intended

Heterogeneity

- Heterogeneous components must be able to interoperate across different:
 - Operating systems
 - Hardware architectures
 - Communication architectures
 - Programming languages
 - Software interfaces
 - Security measures
 - Information representation




Distribution Transparency I

- To hide from the user and the application programmer the separation/distribution of components, so that the system is perceived as a whole rather than a collection of independent components.
- ISO Reference Model for Open Distributed Processing (ODP) identifies the following forms of transparencies:
- Access transparency
 - Access to local or remote resources is identical
 - E.g. Network File System / **Dropbox**
- Location transparency
 - Access without knowledge of location
 - E.g. separation of domain name from machine address.
- Failure transparency
 - Tasks can be completed despite failures
 - E.g. message retransmission, failure of a Web server node should not bring down the website.





Fault Tolerance

- Failure: an offered service no longer complies with its specification (e.g., no longer available or very slow to be usable)
- Fault: cause of a failure (e.g. crash of a component)
- Fault tolerance: no failure despite faults i.e., programmed to handle failures and hides them from users.




Concurrency

- Provide and manage concurrent access to shared resources:
 - Fair scheduling
 - Preserve dependencies (e.g. distributed transactions -- buy a book using Credit card, make sure user has sufficient funds prior to finalizing order)
 - Avoid deadlocks



Openness and Interoperability



- Open system:
"... a system that implements sufficient **open specifications** for interfaces, services, and supporting formats to enable properly engineered applications software to be ported across a wide range of systems with minimal changes, to interoperate with other applications on local and remote systems, and to interact with users in a style which facilitates user portability" (POSIX Open Systems Environment, IEEE POSIX 1003.0)
- Open spec/standard developers - communities:
 - ANSI, IETF, W3C, ISO, IEEE, OMG, Trade associations,...

Summary

- Distributed Systems are everywhere
 - Internet enables users throughout the world to access its (application) services from anywhere
 - Resource sharing is the main motivating factor for constructing distributed systems
 - Construction of DS produces many challenges:
 - Heterogeneity, Openness, Security, Scalability, Failure handling, Concurrency, and Transparency
 - Distributed systems enable globalization:
 - Community (Virtual teams, organizations, social networks)
 - Science (e-Science)
 - Business (..e-Banking..)
 - Entertainment (YouTube, e-Friends)
 - Communication (Zoom,..)
- 
- A colorful illustration of a world map with various cultural icons and the text 'GLOBAL VILLAGE'. The map is green and blue, with icons of people in traditional and modern clothing, a globe, and other symbols representing global connectivity and diversity. The text 'GLOBAL VILLAGE' is written in large, bold, white letters across the center of the map.

